

Author

[Xubo Luo](#)

1. Stable Matching

Gale-Shapley algorithm

```
Gale-Shapley(preference lists of hospitals and students):
  Initialize M to empty matching.
  while(some hospital h is unmatched and has not proposed to every student):
    s <- first student on h's list to whom h has not yet proposed
    if(s is unmatched):
      Add h-s to matching M
    elif(s prefer h to current partner h*):
      Replace h*-s with h-s in matching M
    else:
      s rejects h
  return stable matching M.
```

2. Algorithm Analysis

Big O notation

Upper bounds. $f(n)$ is $O(g(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ such that $0 \leq f(n) \leq c \cdot g(n)$ for all $n \geq n_0$

Big Omega notation

Lower bounds. $f(n)$ is $\Omega(g(n))$ if there exist constants $c > 0$ and $n_0 \geq 0$ such that $f(n) \geq c \cdot g(n) \geq 0$ for all $n \geq n_0$

Big Theta notation

Tight bounds. $f(n)$ is $\Theta(g(n))$ if there exist constants $c_1 > 0$, $c_2 > 0$, and $n_0 \geq 0$ such that $0 \leq c_1 \cdot g(n) \leq f(n) \leq c_2 \cdot g(n)$ for all $n \geq n_0$

3. Graphs

Basic definitions and applications

Trees

Def. An undirected graph is a *tree* if it is connected and does not contain a cycle.

Theorem. Let G be an undirected graph on n nodes. Any two of the following statements imply the third:

- G is connected
- G does not contain a cycle
- G has $n-1$ edges

Graph connectivity and graph traversal

BFS, DFS - $O(m+n)$

- BFS = explore in order of distance from s
- DFS = explore in a different way

Testing bipartiteness

Lemma. If a graph G is bipartite, it cannot contain an odd-length cycle

Lemma. If G be a connected graph, and let L_0, \dots, L_k be the layers produced by BFS starting at node s . Exactly one of the following holds.

1. No edges of G joins two nodes of the same layer, and G is bipartite
2. An edge of G joins two nodes of the same layer, and G contains an odd-length cycle(and hence is not bipartite)

Corollary. A graph G is bipartite **iff** it contains no odd-length cycle.

Connectivity in directed graphs

Def. A graph is *strong connected* if every pair of nodes is mutually reachable.

- Use BFS on G and G^{reverse} to determine if G is strongly connected.

Def. A *strong component* is a maximal subset of mutually reachable nodes.

DAGs and topological ordering

Lemma. G is a DAG **iff** G has a topological ordering.

4.Greedy Algorithm

Coin changing

```
cashier_algorithm(x, c1, c2, ..., cn):
    sort n coin denominations so that  $0 < c1 < c2 < \dots < cn$ 
    S = []
    while(x > 0):
        k = largest coin denomination ck such that ck <= x
        if(no such k):
            return 'no solution'
        else:
            x = x - ck
            S = S + {k}
    return S
```

Interval scheduling

```

earliest finish time first(n, s1, s2, ..., sn, f1, f2, ..., fn):
    sort jobs by finish times and renumber so that  $f_1 \leq f_2 \leq \dots \leq f_n$ 
    S = []
    for j = 1 to n:
        if (job j is compatible with S):
            S = S + {j}
    return S

```

Interval partitioning

- Lecture j starts at s_j and finishes at f_j
- Goal: find minimum number of classrooms to schedule all lectures so that no two lectures occur at the same time in the same room.

```

earliest start time first(n, s1, s2, ..., sn, f1, f2, ..., fn):
    sort lectures by start times and renumber so that  $s_1 \leq s_2 \leq \dots \leq s_n$ 
    d = 0
    for j = 1 to n:
        if (lecture j is compatible with some classroom):
            schedule lecture j in any such classroom k
        else:
            allocate a new classroom d+1
            schedule lecture j in classroom d+1
            d = d+1
    return schedule

```

Scheduling to minimize lateness

- Single resource processes one job at a time
- Job j requires t_j units of processing time and is due at time d_j
- If j starts at time s_j , it finishes at time $f_j = s_j + t_j$
- Lateness: $\alpha_j = \max(0, f_j - d_j)$
- Goal: schedule all jobs to minimize maximum lateness $L = \max_j \alpha_j$

```

Earliest deadline first(n, t1, t2, ..., tn, d1, d2, ..., dn):
    sort jobs by due times and renumber so that  $d_1 \leq d_2 \leq \dots \leq d_n$ 
    t = 0
    for j = 1 to n:
        assign job j to interval[t, t+tj]
        sj = t
        fj = t + tj
        t = t + tj
    return intervals[s1, f1], [s2, f2], ..., [sn, fn]

```

Optimal caching

Caching

- Cache with capacity to store k items
- Sequence of m item requests d_1, d_2, \dots, d_m
- Cache hit: item in cache when requested
- Cache miss: item not in cache when requested

Goal. Eviction schedule that minimize the number of evictions.

Dijkstra's algorithm

For single-source shortest paths problem

```
dijkstra(V,E,len,s):
    foreach v!=s:
        dist[v] = inf
        pred[v] = null
    dist[s] = 0
    foreach v belonging to V:
        insert(pq, v, dist[v])
    while(!empty(pq)):
        u = del-min(pq)
        foreach edge e = (u,v) belonging to E leaving u:
            if dist[v] > dist[u]+len(e):
                decrease-key(pq, v, dist[u]+len(e))
                dist[v] = dist[u]+len(e)
                pred[v] = e
```

Minimum spanning tree

The greedy algorithm

Red rule

- Let C be a cycle with no red edges
- Select an uncolored edge of C of max cost and color it red

Blue rule

- Let D be a cutset with no blue edges
- Select an uncolored edge in D of min cost and color it blue

Greedy algorithm

- Apply the red and blue rules until all edges are colored. The blue edges form an MST
- Note: can stop once n-1 edges colored blue

Prim's algorithm

```
Initialize S = {s} for any node s, T = []
repeat n-1 times:
    Add to T a min-cost edge with exactly one endpoint in S
    Add the other endpoint to S
```

Kruskal's algorithm

```

kruskal(V, E, c):
    sort m edges by cost and renumber so that  $c(e_1) \leq c(e_2) \leq \dots \leq c(e_m)$ 
    T = []
    foreach v belonging to V:
        make-set(v)
    for i=1 to m:
        (u,v) =  $e_i$ 
        if find-set(u) != find-set(v):
            T = T + { $e_i$ }
            union(u,v)
    return T

```

Boruvka's algorithm

Repeat until only one tree.

- Apply blue rule to cutset corresponding to each blue tree
- Color all selected edges blue.

Single-link clustering

k-clustering. Divide objects into *k* non-empty groups

Goal. Given an integer *k*, find a *k*-clustering of maximum spacing.

Min-cost arborescence

Def. Given a digraph $G = (V, E)$ and a root $r \in V$, an **arborescence** (rooted at *r*) is a subgraph $T = (V, F)$ such that

- *T* is a spanning tree of *G* if we ignore the direction of edges.
- There is a (unique) directed path in *T* from *r* to each other node $v \in V$.

```

Edmonds-branching(G, r, c):
    foreach v!=r:
        y(v) = min cost of any edge entering v
         $c'(u,v) = c(u,v) - y(v)$  for each edge (u,v) entering v
    foreach v!=r:
        choose one 0-cost edge entering v and let  $F^*$  be the resulting set of
edges
    if( $F^*$  forms an arborescence):
        return  $T=(V, F^*)$ 
    else:
        C = directed cycle in  $F^*$ 
        contract C to a single supernode, yielding  $G' = (V', E')$ 
         $T' = \text{Edmonds-branching}(G', r, c')$ 
        Extend  $T'$  to an arborescence T in G by adding all but one edge of C
    return T

```

5. Divide and Conquer

Merge sort

```
Merge-Sort(L):
    if (list L has one element):
        return L
    divide the list into two halves A and B
    A = Merge-Sort(A)
    B = Merge-Sort(B)
    L = Merge(A,B)
    return L
```

Counting inversions

```
Sort-And-Count(L):
    if (list L has one element):
        return (0,L)
    Divide the list into two halves A and B
    (ra, A) = Sort-And-Count(A)
    (rb, B) = Sort-And-Count(B)
    (rab, L) = Merge-And-Count(A,B)
    return (ra+rb+rab, L)
```

Randomized quicksort

Goal. Given an array A and pivot element p, partition array so that:

- Smaller element in left subarray L
- Equal elements in middle subarray M
- Larger elements right subarray R

```
Randomized-Quicksort(A):
    if (array A has zero or one element):
        return
    pick pivot p belonging to A uniformly at random
    (L,M,R) = partition-3-way(A,p)
    Randomized-Quicksort(L)
    Randomized-Quicksort(R).
```

Median and selection

```
Quick-Select(A, k):
    pick pivot p belonging to A uniformly at random
    (L,M,R) = Partition-3-way(A,p)
    if (k ≤ |L|):
        return Quick-Select(L,k)
    elif (k > |L|+|M|):
        return Quick-Select(R,k-|L|-|M|)
    elif (k == |L|):
        return p
```

Median-of-median selection algorithm

```
MOM-select(A, k):
```

```

n = |A|
if (n < 50):
    return k-th smallest of element of A via mergesort
group A into n/5 groups of 5 elements each
B = median of each group of 5
p = MOM-select(B, n/10)

(L,M,R) = partition-3-way(A,p)
if (k <= |L|):
    return MOM-select(L,k)
elif (k > |L| + |M|):
    return MOM-select(R, k - |L| - |M|)
else:
    return p

```

Master theorem

Master theorem. Let $a \geq 1$, $b \geq 2$, and $c \geq 0$ and suppose that $T(n)$ is a function on the non-negative integers that satisfies the recurrence

$$T(n) = aT\left(\frac{n}{b}\right) + \Theta(n^c)$$

with $T(0) = 0$ and $T(1) = \Theta(1)$, where n/b means either $\lfloor n/b \rfloor$ or $\lceil n/b \rceil$. Then,

Case 1. If $c > \log_b a$, then $T(n) = \Theta(n^c)$.

Case 2. If $c = \log_b a$, then $T(n) = \Theta(n^c \log n)$.

Case 3. If $c < \log_b a$, then $T(n) = \Theta(n^{\log_b a})$.



Integer multiplication

```

Multiply(x,y,n):
    if (n==1):
        return x * y
    else:
        m = n/2
        a = x/2**m, b = x % 2**m
        c = y/2**m, d = y % 2**m
        e = Multiply(a,c,m)
        f = Multiply(b,d,m)
        g = Multiply(b,c,m)
        h = Multiply(a,d,m)

        return 2^(2m) * e + 2^m * (g+h) + f

```

Matrix multiplication

```

strassen(n,A,B):
    if(n==1):
        return A * B
    partition A and B into sqrt(n)-by-sqrt(n) blocks
    p1 = strassen(n/2, A11, (B11-B22))
    p2 = strassen(n/2, (A11+A12), B22)
    p3 = strassen(n/2, (A21+A22), B11)

```

```

p4 = strassen(n/2, A22, (B21-B11))
p5 = strassen(n/2, (A11+A22), (B11+B22))
p6 = strassen(n/2, (A12-A22), (B21+B22))
p7 = strassen(n/2, (A11-A21), (B11+B12))

c11 = p5+p4-p2+p6
c12 = p1+p2
c21 = p3+p4
c22 = p1+p5-p3-p7
return c

```

Convolution and FFT

```

inverse-FFT(n,y0,y1,...,yn-1):
    if (n==1):
        return y0
    (e0,e1,...,en/2-1) = inverse-FFT(n/2, y0,y2,...,yn-2)
    (d0,d1,...,dn/2-1) = inverse-FFT(n/2, y1,y3,...,yn-1)
    for k = 0 to n/2-1:
        wk = exp(-2*pi*i*k/n)
        ak = ek + w**k * dk
        ak+n/2 = ek - w**k * dk
    return (a0,a1,...,an-1)

```

6.Dynamic Programming

Weight interval scheduling

Greedy method: Earliest finish-time first

```

pj = largest index i<j such that job i is compatible with j.
find-solution(j):
    if (j==0):
        return []
    elif (wi + M[p[j]] > M[j-1]):
        return {j} + find-solution(p[j])
    else:
        return find-solution(j-1)

```

Segmented least squares

$OPT(j)$ = minimum cost for points p_1, p_2, \dots, p_n

e_{ij} = SSE for points p_i, p_{i+1}, \dots, p_j

To compute $OPT(j)$:

- Last segment uses points p_i, \dots, p_j for some $i \leq j$
- Cost = $e_{ij} + c + OPT(i-1)$

Bellman equation:

- $OPT(j) = 0$ if $j=0$
- $OPT(j) = \min_{1 \leq i \leq j} \{e_{ij} + c + OPT(i-1)\}$ if $j > 0$


```

segment-least-squares(n, p1, p2, ..., pn, c):
    for j=1 to n:
        for i=1 to j:
            compute the SSE eij for points pi, ..., pj
    M[0] = 0
    for j=1 to n:
        M[j] = min{eij + c + M[i-1]}
    return M[n]

```

Knapsack problem

```

knapsack(n, w, w1, w1, ..., wn, v1, v2, ..., vn):
    for w=0 to W:
        M[0, w] = 0
    for i=1 to n:
        for w=0 to W:
            if(wi > w):
                M[i, w] = M[i-1, w]
            else:
                M[i, w] = max(M[i-1, w], vi + M[i-1, w-wi])
    return M[n, W]

```

Coin Changing

Given n coin denominations $\{d_1, d_2, \dots, d_n\}$ and a target value V .

Def. $\text{OPT}(v)$ = min number of coins to make change for v

Bellman equation

- $\text{OPT}(v) = 0$, if $v=0$
- $\text{OPT}(v) = \min_{1 \leq i \leq n} \{1 + \text{OPT}(v-d_i)\}$, if $v > 0$

RNA secondary structure

```

rna-secondary-structure(n, b1, b2, ..., bn):
    for k=5 to n-1:
        for i=1 to n-k:
            j = i+k
            compute M[i, j] using formula
    return M[1, n]

```

Sequence alignment

Goal. Given two string x_1, \dots, x_m and y_1, \dots, y_n , find a min-cost alignment

Def. An **alignment** M is a set of ordered pairs x_i-y_j such that each character appears in at most one pair and no crossings.

Def. The **cost** of an alignment M is:

$$\text{cost}(M) = \underbrace{\sum_{(x_i, y_j) \in M} \alpha_{x_i y_j}}_{\text{mismatch}} + \underbrace{\sum_{i: x_i \text{ unmatched}} \delta + \sum_{j: y_j \text{ unmatched}} \delta}_{\text{gap}}$$

```
sequence-alignment(m,n,x1,...,xm,y1,...,yn,theta,alpha):
    for i=0 to m:
        M[i,0] = i*theta
    for j=0 to n:
        M[0,j] = j*theta
    for i=1 to m:
        for j=1 to n:
            M[i,j] = min(alpha[xi][yj] + M[i-1,j-1], theta+M[i-1,j],
            theta+M[i,j-1])
    return M[m,n]
```

Hirschberg's algorithm

Bellman-Ford-Moore algorithm

Lemma. If G has no negative cycle, then there exists a shortest $v \rightarrow t$ path that is simple (and has $\leq n-1$ edges)

Def. $\text{OPT}(i,v)$ = length of shortest $v \rightarrow t$ path that uses $\leq i$ edges

Goal. $\text{OPT}(n-1, v)$ for each v

Bellman equation.

- $\text{OPT}(i,v) = 0$, if $i = 0$ and $v = t$
- $\text{OPT}(i,v) = \text{inf}$, if $i = 0$ and $v \neq t$
- $\text{OPT}(i,v) = \min(\text{OPT}(i-1,v), \min\{\text{OPT}(i-1,w) + \text{len}(v,w)\})$, if $i > 0$

```
shortest-paths(V,E,len,t):
    foreach node v belonging to V:
        M[0,v] = inf
    M[0,t] = 0
    for i = 1 to n-1:
        foreach node v belonging to V:
            M[i,v] = M[i-1,v]
            foreach edge(v,w) belonging to E:
                M[i,v] = min{M[i,v], M[i-1,w] + len(v,w)}
```

Space optimization.

Maintain two 1D arrays

- $d[v]$ = length of a shortest $v \rightarrow t$ path that we have found so far
- $\text{successor}[v]$ = next node on a $v \rightarrow t$ path

```

Bellman-Ford-Moore(V, E, c, t):
    foreach node v belonging to V:
        d[v] = inf
        successor[v] = null
    d[t] = 0
    for i = 1 to n-1:
        foreach node w belonging to V:
            if(d[v] > d[w] + len(v,w)):
                d[v] = d[w] + len(v,w)
                successor[v] = w
    if (no d[] value changed in pass i):
        stop

```

Distance-vector protocols

Communication network

- Node = router
- Edge = direct communication link
- Length of edge = latency of link

Negative cycles

Negative cycle detection problem

Given a digraph $G = (V, E)$, with edge length $\text{len}(v, w)$, find a negative cycle (if one exists)

- Run Bell-Ford-Moore on G' for $n'=n+1$ passes (instead of $n'-1$)
- If no $d[v]$ values updated in pass n' , then no negative cycles
- Otherwise, suppose $d[s]$ updated in pass n'
- Define $\text{pass}(v)$ = last pass in which $d[v]$ was updated
- Observe $\text{pass}(s) = n'$ and $\text{pass}(\text{successor}[v]) \geq \text{pass}(v) - 1$ for each v
- Following successor pointers, we must eventually repeat a node.
- The corresponding cycle is a negative cycle`

7. Network Flow

max-flow and min-cut problems

Def. A cut's **capacity** is the sum of the capacities of the edges from A to B

$$\text{cap}(A, B) = \sum_{e \text{ out of } A} c(e)$$

Def. The **value** of a flow f is: $\text{val}(f) = \sum_{e \text{ out of } s} f(e) - \sum_{e \text{ in to } s} f(e)$

Ford-Fulkerson algorithm

Def. An augmenting path is a simple $s \rightarrow t$ path in the residual network G_f

```

Augment(f,c,P):
    theta = bottleneck capacity of augment path P
    foreach edge e belonging to P:
        if (e belongs to E):
            f(e) = f(e) + theta
        else:
            f(e_reverse) = f(e_reverse) - theta
    return f

```

```

ford-fulkerson(G):
    foreach edge e belonging to E:
        f(e) = 0
    while (there exists an s->t path P in Gf):
        f = Augment(f,c,P)
        update Gf
    return f

```

max-flow min-cut theorem

$\text{val}(f) \leq \text{cap}(A,B)$

if $\text{val}(f) = \text{cap}(A,B)$: then f is a max flow and (A,B) is a min cut

Augmenting path theorem. A flow f is a max flow **iff** no augmenting paths.

capacity-scaling algorithm

Goal. Choose augmenting paths so that:

- Can find augmenting paths efficiently
- Few iterations

```

capacity-scaling(G):
    foreach edge e belonging to E:
        f(e) = 0
    delta = largest power of 2 <= C
    while (delta >= 1):
        Gf(delta) = delta-residual network of G with respect to flow f
        while (there exists an s->t path P in Gf(delta)):
            f = Augment(f,c,P)
            update Gf(delta)
        delta = delta/2
    return f

```

shortest augmenting paths

Q: How to choose next augmenting path in Ford-Fulkerson?

A: Pick one that uses the fewest edges.

```

short-augmenting-path(G):
    foreach e belonging to E:
        f(e) = 0
        Gf = residual network of G with respect to flow f
        while (there exists an s->t path in Gf):
            P = BFS(Gf)
            f = Augment(f, c, P)
            update Gf
    return f

```

Dinitz algorithm

Two types of augmentations:

- Normal: length of shortest path does not change
- Special: length of shortest path strictly increases

```

initialize(G, f):
    LG = level-graph of Gf
    P = []
    GOTO advance(s)

advance(v):
    if (v == t):
        augment(P)
        remove saturated edges from LG
        P = []
        GOTO advance(s)

repeat(v):
    if (v == s):
        stop
    else:
        delete v (and all incident edges) from LG
        remove last edge(u, v) from P
        GOTO advance(u)

```

simple unit-capacity network

Def. A flow network is a **simple unit-capacity network** if:

- Every edge has capacity 1
- Every node (other than s or t) has exactly one entering edge, or exactly one leaving edge, or both

Bipartite matching

Given a bipartite graph $G = (L \cup R, E)$, find a max-cardinality matching

Formulation

- Create digraph $G' = \{L \cup R \cup \{s, t\}, E'\}$
- Direct all edges from L to R, and assign infinite (or unit) capacity
- Add unit-capacity edges from s to each node in L
- Add unit-capacity edges from each node in R to t

Disjoint paths

Edge-disjoint paths problem

Given a digraph $G=(V,E)$ and two nodes s and t , find the max number of edge-disjoint $s \rightarrow t$ paths

Menger's theorem The max number of edge-disjoint $s \rightarrow t$ paths equals the min number of edges whose removal disconnects t from s .

Max-flow formulation. Assign unit capacity to every edge

Edge-disjoint paths problem in undirected graphs

Given a graph $G=(V,E)$ and two nodes s and t , find the max number of edge-disjoint $s-t$ paths

Max-flow formulation. Replace each edge with two antiparallel edges and assign unit capacity to every edge

Theorem. Max number of edge-disjoint $s \rightarrow t$ paths = value of max flow.

Extensions to max flow

Multiple sources and sinks: max-flow formulation

- Add a new source node s and sink node t
- For each original source node s_i , add edge (s, s_i) with capacity inf
- For each original sink node t_j , add edge (t_j, t) with capacity inf

Circulation with supplies and demands

Def Given a digraph $G = (V, E)$ with edge capacities $c(e) \geq 0$ and node demands $d(v)$, a circulation is a function $f(e)$ that satisfies:

- For each e belonging to E : $0 \leq f(e) \leq c(e)$
- For each v belonging to V : $\sum_{\text{in}} f(e) - \sum_{\text{out}} f(e) = d(v)$
- Add new source s and sink t
- For each v with $d(v) < 0$, add edge (s, v) with capacity $-d(v)$
- For each v with $d(v) > 0$, add edge (v, t) with capacity $d(v)$

Claim:

G has a circulation **iff** G' has a max flow of value $D = \sum_{\text{positive}} d(v) = \sum_{\text{negative}} -d(v)$

Survey design

- Design survey asking n_1 consumers about n_2 products.
- Can survey consumer i about product j only if they own it.
- Ask consumer i between c_i and c_i' questions.
- Ask between p_j and p_j' consumers about product j .

Goal. Design a survey that meets these specs, if possible

Max-flow formulation. Model as a circulation problem with lower bounds.

- Add edge (i, j) if consumer j owns product i .
- Add edge from s to consumer j .
- Add edge from product i to t .
- Add edge from t to s .
- All demands = 0.

- Integer circulation \Leftrightarrow feasible survey design

Airline scheduling

- Manage flight crews by reusing them over multiple flights.
- Input: set of k flights for a given day.
- Flight i leaves origin o_i at time s_i and arrives at destination d_i at time f_i .
- Minimize number of flight crews.

Circulation formulation. [to see if c crews suffice]

- For each flight i , include two nodes u_i and v_i .
- Add source s with demand $-c$, and edges (s, u_i) with capacity 1.
- Add sink t with demand c , and edges (v_i, t) with capacity 1.
- For each i , add edge (u_i, v_i) with lower bound and capacity 1.
- if flight j reachable from i , add edge (v_i, u_j) with capacity 1

Image segmentation

Projection selection

Projects with prerequisites.

- Set of possible projects P : project v has associated revenue p_v .
- Set of prerequisites E : $(v, w) \in E$ means w is a prerequisite for v .
- A subset of projects $A \subseteq P$ is feasible if the prerequisite of every project in A also belongs to A .

Project selection problem. Given a set of projects P and prerequisites E , choose a feasible subset of projects to maximize revenue.

Min-cut formulation

- Assign a capacity of ∞ to each prerequisite edge.
- Add edge (s, v) with capacity p_v if $p_v > 0$.
- Add edge (v, t) with capacity $-p_v$ if $p_v < 0$.
- For notational convenience, define $p_s = p_t = 0$.

Claim. (A, B) is min cut $\iff A - \{s\}$ is an optimal set of projects.

Assignment problem

Input. Weighted, complete bipartite graph $G = (X \cup Y, E)$ with $|X| = |Y|$.

Goal. Find a perfect matching of min weight.

Bipartite matching. Can solve via reduction to maximum flow.

Input-queued switching

Input-queued switch

- n input ports and n output ports in an n -by- n crossbar layout.
- At most one cell can depart an input at a time.
- At most one cell can arrive at an output at a time.
- Cell arrives at input x and must be routed to output y .

8. Intractability

Poly-time reductions

A working definition. Those with poly-time algorithms.

yes	probably no
shortest path	longest path
min cut	max cut
2-satisfiability	3-satisfiability
planar 4-colorability	planar 3-colorability
bipartite vertex cover	vertex cover
matching	3d-matching
primality testing	factoring
linear programming	integer linear programming

Reduction. Problem X *polynomial-time reduces to* problem Y if arbitrary instances of problem X can be solved using:

- Polynomial number of standard computational steps, plus
- Polynomial number of calls to oracle that solves problem Y

Notation. $X \leq_p Y$

If x reduces to y , then x is easier than y .

Establish equivalence. If both $x \leq_p y$ and $y \leq_p x$, we use notation $x =_p y$.

Packing and covering problems

Independent-Set. Given a graph $G = (V, E)$ and an integer k , is there a subset of k (or more) vertices such that no two are adjacent?

Vertex-Cover. Given a graph $G = (V, E)$ and an integer k , is there a subset of k (or fewer) vertices such that each edge is incident to at least one vertex in the subset?

Theorem. $\text{INDEPENDENT-SET} \equiv_p \text{VERTEX-COVER}$.

Set-Cover. Given a set U of elements, a collection S of subsets of U , and an integer k , are there $\leq k$ of these subsets whose union is equal to U ?

Theorem. $\text{Vertex-Cover} \leq_p \text{Set-Cover}$

Lemma. $G = (V, E)$ contains a vertex cover of size k iff (U, S, k) contains a set cover of size k .

Constraint satisfaction problems

Satisfiability

Literal. A Boolean variable or its negation

Clause. A disjunction of literals.

Conjunctive normal form (CNF). A propositional formula Φ that is a conjunction of clauses.

SAT. Given a CNF formula Φ , does it have a satisfying truth assignment?

3-SAT. SAT where each clause contains exactly 3 literals (and each literal corresponds to a different variable).

Theorem. $3\text{-SAT} \leq_p \text{independent-set}$

Construction.

- G contains 3 nodes for each clause, one for each literal.
- Connect 3 literals in a clause in a triangle.
- Connect literal to each of its negations.

Lemma. Φ is satisfiable **iff** G contains an independent set of size $k = |\Phi|$

Basic reduction strategies

- Simple equivalence: $\text{INDEPENDENT-SET} \equiv_p \text{VERTEX-COVER}$.
- Special case to general case: $\text{VERTEX-COVER} \leq_p \text{SET-COVER}$.
- Encoding with gadgets: $3\text{-SAT} \leq_p \text{INDEPENDENT-SET}$.

Decision problem. Does there exist a vertex cover of size $\leq k$?

Search problem. Find a vertex cover of size $\leq k$

Optimization problem. Find a vertex cover of minimum size

$\text{Vertex-Cover} =_p \text{Find-Vertex-Cover} =_p \text{Find-Min-Vertex-Cover}$

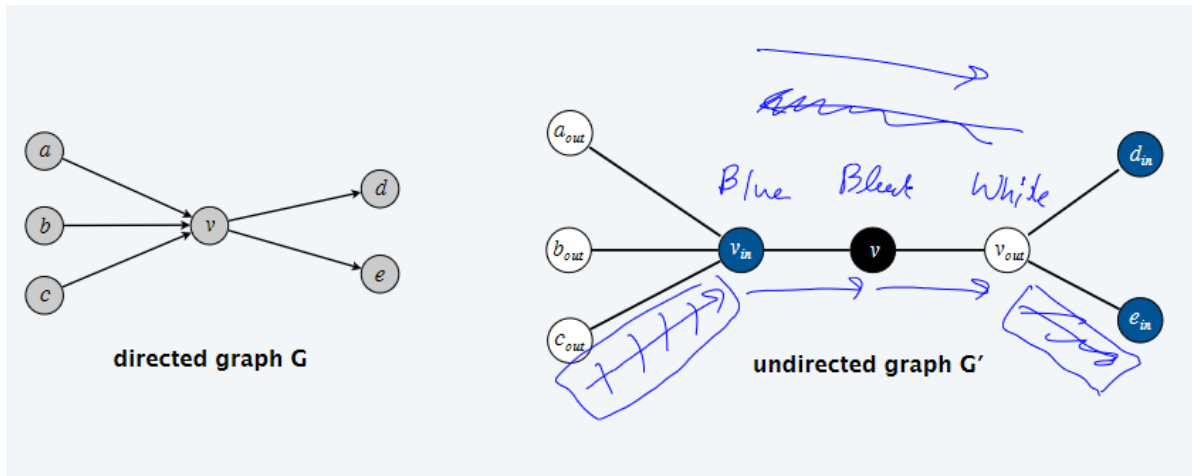
Sequencing problems

HAMILTON-CYCLE. Given an undirected graph $G = (V, E)$, does there exist a cycle Γ that visits every node exactly once?

DIRECTED-HAMILTON-CYCLE. Given a directed graph $G = (V, E)$, does there exist a directed cycle Γ that visits every node exactly once?

Theorem. $\text{DIRECTED-HAMILTON-CYCLE} \leq_p \text{HAMILTON-CYCLE}$.

Pf.

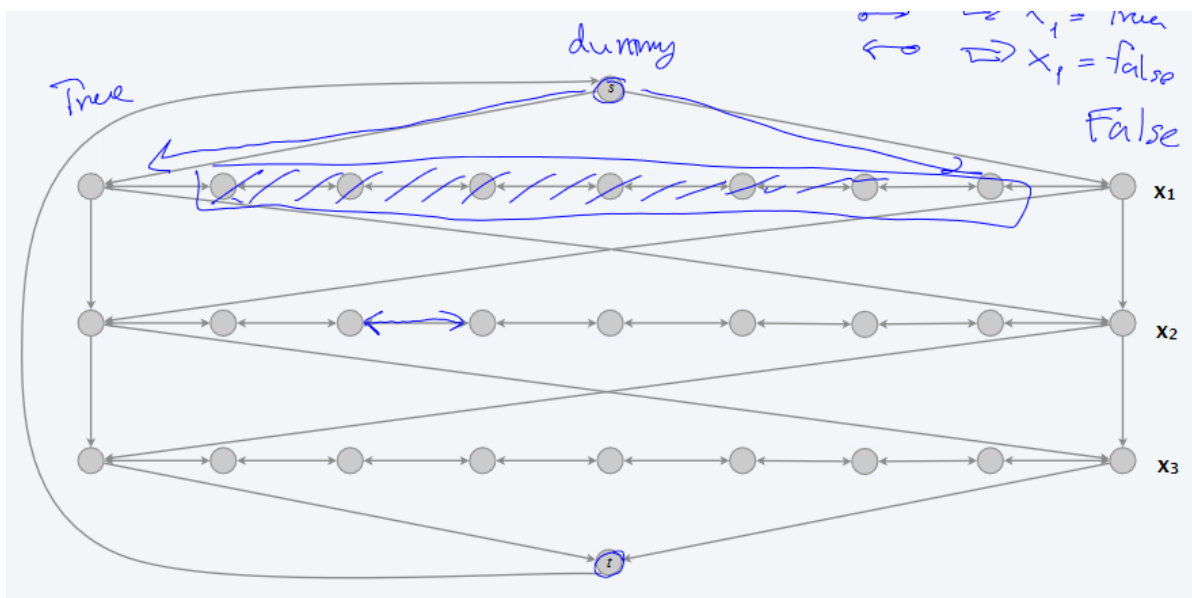


Lemma. G has a directed Hamilton cycle **iff** G' has a Hamilton cycle

Theorem. $3\text{-SAT} \leq_p \text{DIRECTED-HAMILTON-CYCLE}$.

Construction. Given 3-SAT instance Φ with n variables x_i and k clauses.

- Construct G to have 2^n Hamilton cycles.
- Intuition: traverse path i from left to right \Leftrightarrow set variable $x_i = \text{true}$



Construction. Given 3-SAT instance Φ with n variables x_i and k clauses.

- For each clause: add a node and 2 edges per literal.

Lemma. Φ is satisfiable **iff** G has a Hamilton cycle.

Partitioning problems

3-dimensional matching

3D-Matching. Given 3 disjoint sets X , Y , and Z , each of size n and a set $T \subseteq X \times Y \times Z$ of triples, does there exist a set of n triples in T such that each element of $X \cup Y \cup Z$ is in exactly one of these triples?

Theorem. $3\text{-SAT} \leq_p 3\text{D-MATCHING}$

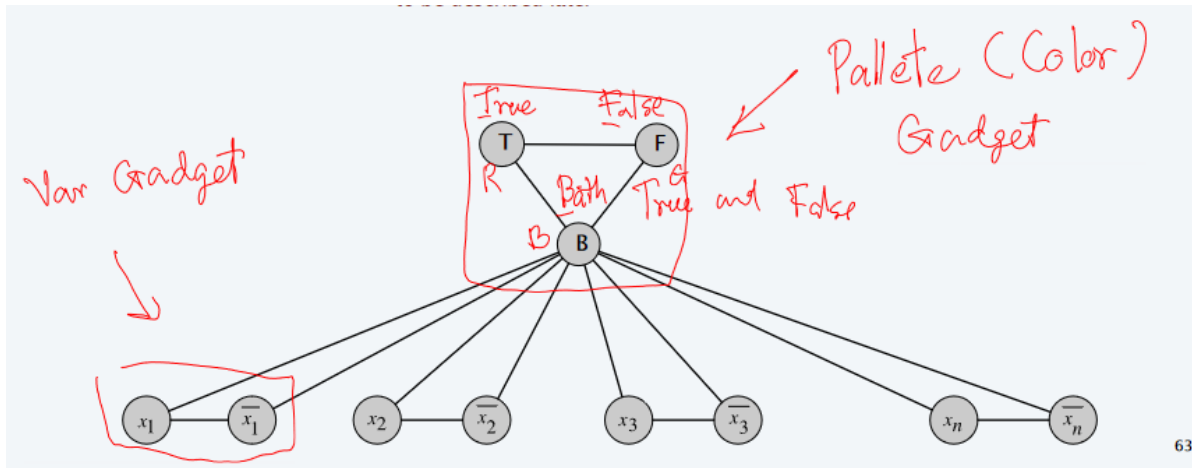
Graph coloring

3-COLOR. Given an undirected graph G , can the nodes be colored black, white, and blue so that no adjacent nodes have the same color?

Theorem. $3\text{-SAT} \leq_P 3\text{-COLOR}$.

Construction

- Create a graph G with a node for each literal.
- Connect each literal to its negation.
- Create 3 new nodes T , F , and B ; connect them in a triangle.
- Connect each literal to B .
- For each clause C_j , add a gadget of 6 nodes and 13 edges.



63

Numerical problems

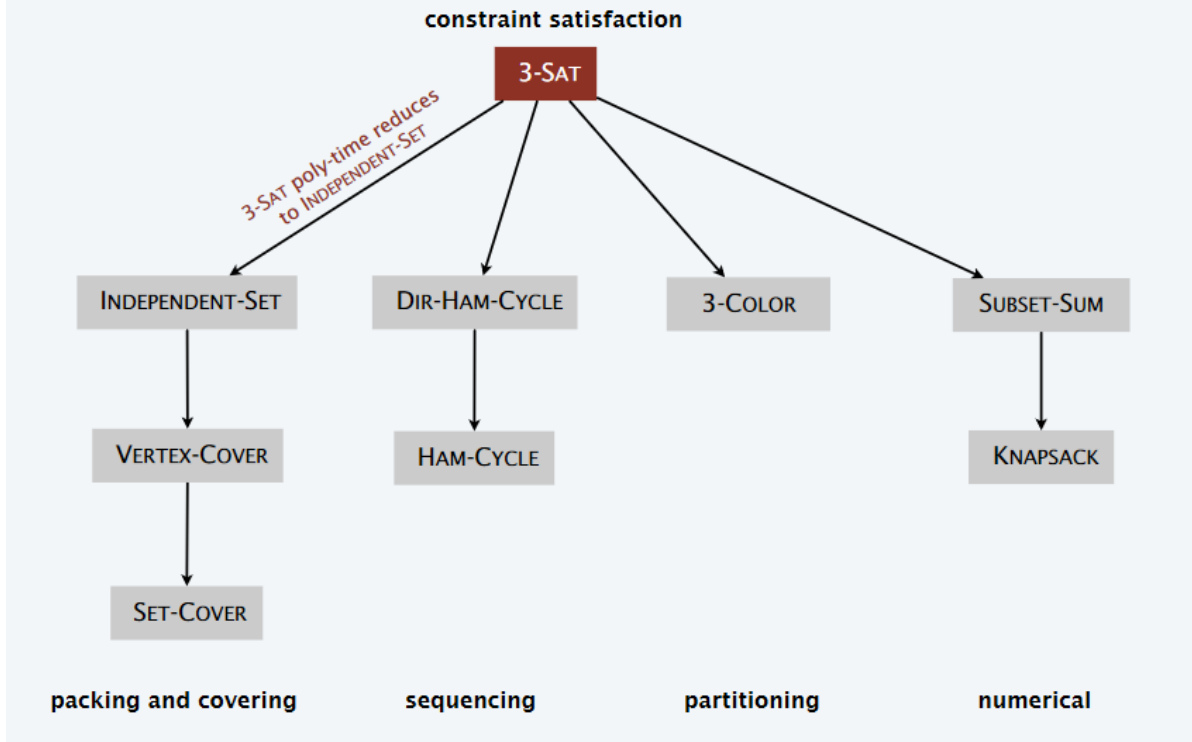
Subset sum

Subset sum. Given n natural numbers w_1, \dots, w_n and an integer W , is there a subset that adds up to exactly W ?

Theorem. $3\text{-SAT} \leq_P \text{SUBSET-SUM}$.

Poly-time reduction

Poly-time reductions



P vs. NP

P

Decision problem

- Problem X is a set of strings
- Instance s is one string
- Algorithm A solves problem X : $A(s) = \text{yes}$ if s belongs to X ; no if s doesn't belong to X

Def. Algorithm A runs in **polynomial time** if for every string s , $A(s)$ terminates in $\leq p(|s|)$ "steps," where $p(\cdot)$ is some polynomial function.

Def. P = set of decision problems for which there exists a poly-time algorithm

NP

Def. Algorithm $C(s, t)$ is a **certifier** for problem X if for every string $s : s \in X$ **iff** there exists a string t such that $C(s, t) = \text{yes}$.

Def. **NP** = set of decision problems for which there exists a poly-time certifier.

- $C(s, t)$ is a poly-time algorithm.
- Certificate t is of polynomial size: $|t| \leq p(|s|)$ for some polynomial $p(\cdot)$

P. Decision problems for which there exists a poly-time algorithm.

NP. Decision problems for which there exists a poly-time certifier.

EXP. Decision problems for which there exists an exponential-time algorithm.

NP-complete

NP-complete. A problem $Y \in \text{NP}$ with the property that for every problem $X \in \text{NP}$, $X \leq_p Y$.

Theorem. $\text{SAT} \in \text{NP-complete}$

Recipe. To prove that $Y \in \text{NP-complete}$:

- Step 1. Show that $Y \in \text{NP}$.
- Step 2. Choose an NP-complete problem X .
- Step 3. Prove that $X \leq_p Y$.

Proposition. If $X \in \text{NP-complete}$, $Y \in \text{NP}$, and $X \leq_p Y$, then $Y \in \text{NP-complete}$.

co-NP

Def. Given a decision problem X , its complement $\text{not}(X)$ is the same problem with the yes and no answers reversed.

co-NP. Complements of decision problems in NP.

Theorem. If $\text{NP} \neq \text{co-NP}$, then $\text{P} \neq \text{NP}$.

NP-hard

NP-hard. A problem such that every problem in NP poly-time reduces to it

Summarize

- P: 多项式时间内可解的问题
- NP: 多项式时间内可验证的问题
- NPC: 所有NP问题都可以归约到这个NP问题
- NP-Hard: 所有NP问题都可以归约到这个问题

